

Avaliacao_ensemble_test

November 4, 2019

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
from datetime import datetime
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.model_selection import StratifiedKFold, GridSearchCV
from sklearn.ensemble import RandomForestClassifier,
↳ GradientBoostingClassifier, AdaBoostClassifier, VotingClassifier
from sklearn.metrics import roc_auc_score, roc_curve, auc,
↳ precision_recall_curve
from sklearn.metrics import classification_report, confusion_matrix
from xgboost import XGBClassifier
from mlxtend.plotting import plot_learning_curves
from yellowbrick.model_selection import LearningCurve
import matplotlib.gridspec as gridspec
import itertools
from sklearn.model_selection import cross_val_score, train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn import tree
from sklearn.naive_bayes import MultinomialNB
from sklearn.utils import shuffle
from sklearn.preprocessing import LabelEncoder, OrdinalEncoder

[2]: def timer(start_time=None):
    if not start_time:
        start_time = datetime.now()
        return start_time
    elif start_time:
        tmin, tsec = divmod((datetime.now() - start_time).total_seconds(), 60)
        print('\n Tempo Necessário: %i minutos and %s segundos.' % (tmin,
↳ round(tsec, 2)))

[3]: train = pd.read_csv('trainLR.csv')
train = shuffle(train)
```

```
X_train = train.iloc[:,1:25]
Y_train = train.loc[:, train.columns == 'Y']
test = pd.read_csv('testLR.csv')
test = shuffle(test)
X_test = test.iloc[:,1:25]
Y_test = test.loc[:, test.columns == 'Y']
```

```
[4]: print(X_train.shape)
```

```
(109992, 24)
```

```
[5]: print(X_test.shape)
```

```
(65995, 24)
```

```
[6]: X_train.head()
```

```
[6]:
```

	v1	v10	v29	v87	v111	v277	v279	v280	v281	v282	\
95975	0	0	0	0	0	0.758621	0.0	0.866667	0.740741	0.636364	
21840	0	0	0	0	0	0.000000	-1.0	0.000000	0.000000	-1.000000	
68686	0	0	0	0	0	1.000000	-1.0	1.000000	1.000000	-1.000000	
107302	0	0	0	0	0	0.857143	-1.0	0.857143	0.857143	-1.000000	
1795	0	0	0	2	2	0.958333	0.0	1.000000	0.950000	1.000000	

	...	v605	v606	v607	v608	v609	v610	v681	v683	v684	\
95975	...	0.833333	0.833333	0.866667	0.7	0.7	0.7	1	0	0	
21840	...	0.000000	0.000000	0.000000	-1.0	-1.0	-1.0	0	0	0	
68686	...	1.000000	1.000000	1.000000	-1.0	-1.0	-1.0	0	0	0	
107302	...	1.000000	1.000000	1.000000	-1.0	-1.0	-1.0	0	0	0	
1795	...	0.950000	0.950000	0.950000	1.0	1.0	1.0	0	0	0	

	v691
95975	0.400000
21840	0.500000
68686	0.500000
107302	0.363636
1795	0.000000

```
[5 rows x 24 columns]
```

```
[7]: # Fit a Decision Tree model as comparison
starttime = timer(None)
start_time = timer(None)
clf_D DecisionTreeClassifier = DecisionTreeClassifier()
clf_D DecisionTreeClassifier.fit(X_train, Y_train.values.ravel())
DecisionTreeClassifier_pred = clf_D DecisionTreeClassifier.predict(X_test)
timer(start_time)
```

```

accuracy_score(Y_test, DecisionTreeClassifier_pred)
false_positive_rate, true_positive_rate, thresholds = roc_curve(Y_test,
    ↳ DecisionTreeClassifier_pred)
roc_auc = auc(false_positive_rate, true_positive_rate)

plt.figure(1)
matrix_DecisionTreeClassifier = confusion_matrix(Y_test,
    ↳ DecisionTreeClassifier_pred)
plt.figure(figsize=(9,5))
DecisionTreeClassifier = sns.heatmap(matrix_DecisionTreeClassifier, annot=True,
    ↳ cbar=False, fmt="d", cmap='coolwarm', linecolor='black', linewidths = 2)
bottom, top = DecisionTreeClassifier.get_yylim()
DecisionTreeClassifier.set_yylim(bottom + 0.5, top - 0.5)
plt.ylabel('Rótulos Verdadeiro')
plt.xlabel('Predicted Label')
plt.title('Matriz de Confusão')
plt.show()

plt.figure(2)
plt.figure(figsize=(9,5))
plt.title('Receiver Operating Characteristic')
plt.plot(false_positive_rate, true_positive_rate, 'b',
    label='AUC = %0.2f'% roc_auc)
plt.legend(loc='lower right')
plt.plot([0,1],[0,1], 'r--')
plt.xlim([-0.0,1.0])
plt.ylim([-0.0,1.0])
plt.ylabel('Taxa de verdadeiros Positivos')
plt.xlabel('Taxa de Falsos Positivos')
plt.show()


print("Relatório de Classificação")
print(classification_report(Y_test, DecisionTreeClassifier_pred))

print("Acurácia do Modelo")
accuracy_score(Y_test, DecisionTreeClassifier_pred)


```

Tempo Necessário: 0 minutos and 0.73 segundos.

<Figure size 432x288 with 0 Axes>



<Figure size 432x288 with 0 Axes>



Relatório de Classificação

	precision	recall	f1-score	support
0	0.31	0.23	0.26	13442
1	0.82	0.87	0.84	52553
accuracy			0.74	65995
macro avg	0.56	0.55	0.55	65995
weighted avg	0.71	0.74	0.72	65995

Acurácia do Modelo

[7]: 0.739631790287143

Fit a Simple Random Forest model

```
[8]: starttime = timer(None)
start_time = timer(None)
clf = RandomForestClassifier(n_estimators=100,
    ↳max_features="auto",random_state=0)
clf.fit(X_train, Y_train.values.ravel())
scores = cross_val_score(clf, X_train, Y_train.values.ravel(), cv=3)
RandomForestClassifier_pred = clf.predict(X_test)
timer(start_time)
accuracy_score(Y_test, RandomForestClassifier_pred)
false_positive_rate, true_positive_rate, thresholds = roc_curve(Y_test,
    ↳RandomForestClassifier_pred)
roc_auc = auc(false_positive_rate, true_positive_rate)

plt.figure(1)
matrix_RandomForestClassifier = confusion_matrix(Y_test,
    ↳RandomForestClassifier_pred)
plt.figure(figsize=(9,5))
map_RandomForestClassifier = sns.heatmap(matrix_RandomForestClassifier,
    ↳annot=True, cbar=False, fmt="d", cmap='coolwarm', linecolor='black',
    ↳linewidths = 1)
bottom, top = map_RandomForestClassifier.get_ylim()
map_RandomForestClassifier.set_ylim(bottom + 0.5, top - 0.5)
plt.ylabel('Rótulos Verdadeiro')
plt.xlabel('Predicted Label')
plt.title('Matriz de Confusão')
plt.show()

plt.figure(2)
plt.figure(figsize=(9,5))
plt.title('Receiver Operating Characteristic')
plt.plot(false_positive_rate, true_positive_rate, 'b',
```

```

label='AUC = %0.2f'% roc_auc)
plt.legend(loc='lower right')
plt.plot([0,1],[0,1], 'r--')
plt.xlim([-0.0,1.0])
plt.ylim([-0.0,1.0])
plt.ylabel('Taxa de verdadeiros Positivos')
plt.xlabel('Taxa de Falsos Positivos')
plt.show()

print("Relatório de Classificação")
print(classification_report(Y_test, RandomForestClassifier_pred))


print("Acurácia do Modelo")
accuracy_score(Y_test, RandomForestClassifier_pred)

print("Acurácia do Modelo Cross Validation")
print(scores.mean())


```

Tempo Necessário: 0 minutos and 36.27 segundos.

<Figure size 432x288 with 0 Axes>



<Figure size 432x288 with 0 Axes>



Relatório de Classificação

	precision	recall	f1-score	support
0	0.41	0.14	0.21	13442
1	0.81	0.95	0.88	52553
accuracy			0.78	65995
macro avg	0.61	0.54	0.54	65995
weighted avg	0.73	0.78	0.74	65995

Acurácia do Modelo

Acurácia do Modelo Cross Validation

0.78217505526712

Fit a AdaBoost model

```
[9]: starttime = timer(None)
start_time = timer(None)
clf = AdaBoostClassifier(n_estimators=100)
clf.fit(X_train, Y_train.values.ravel())
AdaBoostClassifier_pred = clf.predict(X_test)
scores = cross_val_score(clf, X_train, Y_train.values.ravel(), cv=3)
timer(start_time)
accuracy_score(Y_test, AdaBoostClassifier_pred)
```

```

false_positive_rate, true_positive_rate, thresholds = roc_curve(Y_test,
↳AdaBoostClassifier_pred)
roc_auc = auc(false_positive_rate, true_positive_rate)

plt.figure(1)
matrix_AdaBoostClassifier = confusion_matrix(Y_test, AdaBoostClassifier_pred)
plt.figure(figsize=(9,5))
map_matrix_AdaBoostClassifier = sns.heatmap(matrix_AdaBoostClassifier,
↳annot=True, cbar=False, fmt="d", cmap = 'coolwarm', linecolor = 'black',
↳linewidths = 1)
bottom, top = map_matrix_AdaBoostClassifier.get_ylim()
map_matrix_AdaBoostClassifier.set_ylim(bottom + 0.5, top - 0.5)
plt.ylabel('Rótulos Verdadeiro')
plt.xlabel('Predicted Label')
plt.title('Matriz de Confusão')
plt.show()

plt.figure(2)
plt.figure(figsize=(9,5))
plt.title('Receiver Operating Characteristic')
plt.plot(false_positive_rate, true_positive_rate, 'b',
label='AUC = %0.2f'% roc_auc)
plt.legend(loc='lower right')
plt.plot([0,1],[0,1], 'r--')
plt.xlim([-0.0,1.0])
plt.ylim([-0.0,1.0])
plt.ylabel('Taxa de verdadeiros Positivos')
plt.xlabel('Taxa de Falsos Positivos')
plt.show()

print("Relatório de Classificação")
print(classification_report(Y_test, AdaBoostClassifier_pred))


print("Acurácia do Modelo")
accuracy_score(Y_test, AdaBoostClassifier_pred)

print("Acurácia do Modelo Cross Validation")
print(scores.mean())


```

Tempo Necessário: 0 minutos and 20.25 segundos.

<Figure size 432x288 with 0 Axes>



<Figure size 432x288 with 0 Axes>



Relatório de Classificação

	precision	recall	f1-score	support
0	0.61	0.08	0.14	13442
1	0.81	0.99	0.89	52553
accuracy			0.80	65995
macro avg	0.71	0.53	0.52	65995
weighted avg	0.77	0.80	0.74	65995

Acurácia do Modelo

Acurácia do Modelo Cross Validation

0.7991672041264524

Fit a Gradient Boosting model

```
[10]: gb_clf2 = GradientBoostingClassifier(n_estimators=200, learning_rate=0.5,
    ↳max_features=2, max_depth=2, random_state=1)
gb_clf2.fit(X_train, Y_train.values.ravel())
GradientBoostingClassifier_predictions = gb_clf2.predict(X_test)

matrix_GradientBoostingClassifier2 = confusion_matrix(Y_test,
    ↳GradientBoostingClassifier_predictions)
scores = cross_val_score(gb_clf2, X_train, Y_train.values.ravel(), cv=3)
plt.figure(figsize=(9,5))
map_matrix_GradientBoostingClassifier2 = sns.
    ↳heatmap(matrix_GradientBoostingClassifier2, annot=True, cbar=False, fmt="d",
    ↳cmap = 'coolwarm', linecolor = 'black', linewidths = 1)
bottom, top = map_matrix_GradientBoostingClassifier2.get_ylim()
map_matrix_GradientBoostingClassifier2.set_ylim(bottom + 0.5, top - 0.5)
plt.ylabel('Rótulos Verdadeiro')
plt.xlabel('Predicted Label')
plt.title('Matriz de Confusão')

print("Relatório de Classificação")
print(classification_report(Y_test, GradientBoostingClassifier_predictions))

print("Acurácia do Modelo")
accuracy_score(Y_test, GradientBoostingClassifier_predictions)

print("Acurácia do Modelo Cross Validation")
print(scores.mean())
```

Relatório de Classificação


	precision	recall	f1-score	support
0	0.59	0.10	0.18	13442

1	0.81	0.98	0.89	52553
accuracy			0.80	65995
macro avg	0.70	0.54	0.53	65995
weighted avg	0.77	0.80	0.74	65995

Acurácia do Modelo

Acurácia do Modelo Cross Validation

0.8000308936340507



XGboost Classifier

```
[11]: starttime = timer(None)
start_time = timer(None)
xgb_clf = XGBClassifier(base_score=0.2, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=0.1, colsample_bytree=1, gamma=1,
                        learning_rate=0.1, max_delta_step=1, max_depth=3,
                        min_child_weight=1, missing=None, n_estimators=300, n_jobs=4,
                        nthread=None, objective='reg:squarederror', random_state=1,
                        reg_alpha=0, reg_lambda=10, scale_pos_weight=1, seed=None,
                        silent=None, subsample=1, verbosity=1)
xgb_clf.fit(X_train, Y_train.values.ravel())
predictions_xgb = xgb_clf.predict(X_test)
scores = cross_val_score(xgb_clf, X_train, Y_train.values.ravel(), cv=3)
timer(start_time)
```

```

false_positive_rate, true_positive_rate, thresholds = roc_curve(Y_test,
    ↪ predictions_xgb)
roc_auc = auc(false_positive_rate, true_positive_rate)
matrix_xgb_clf = confusion_matrix(Y_test, predictions_xgb)

plt.figure(1)
plt.figure(figsize=(9,5))
xgb_clf_heatmap = sns.heatmap(matrix_xgb_clf,annot=True, cbar=False, fmt="d",
    ↪ cmap = 'coolwarm', linecolor = 'black', linewidths = 1)
bottom, top = xgb_clf_heatmap.get_ylim()
xgb_clf_heatmap.set_ylim(bottom + 0.5, top - 0.5)
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.title('Confusion Matrix')
plt.show()

plt.figure(2)
plt.figure(figsize=(9,5))
plt.title('Receiver Operating Characteristic')
plt.plot(false_positive_rate, true_positive_rate, 'b',
label='AUC = %0.2f'% roc_auc)
plt.legend(loc='lower right')
plt.plot([0,1],[0,1], 'r--')
plt.xlim([-0.0,1.0])
plt.ylim([-0.0,1.0])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()

plt.figure(3)
precision, recall, thresholds = precision_recall_curve(Y_test, predictions_xgb)
plt.figure(figsize = (9,5))
plt.plot(recall, precision)
plt.plot([0, 1], [0.5, 0.5], linestyle = '--')
plt.xlabel('Recall', fontsize = 16)
plt.ylabel('Precision', fontsize = 16)
plt.xticks(size = 18)
plt.yticks(size = 18)
plt.title('Precision-Recall', fontsize = 28)
plt.show();

print("Classification Report")
print(classification_report(Y_test, predictions_xgb))

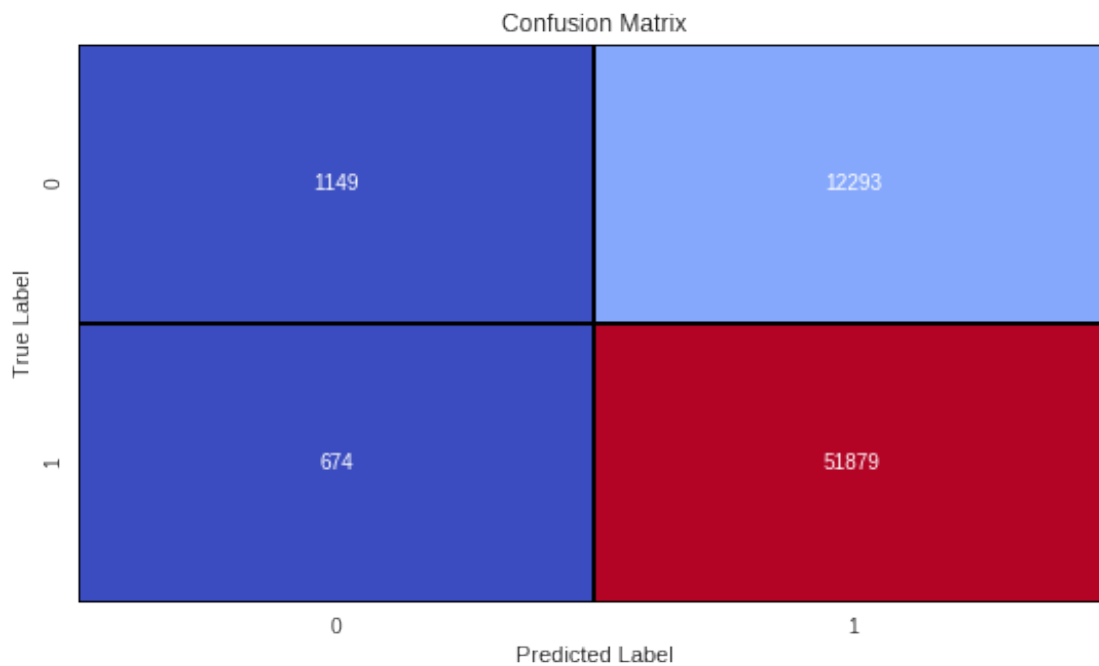
```

```
print("Acurácia do Modelo")
accuracy_score(Y_test, predictions_xgb)


print("Acurácia do Modelo Cross Validation")
print(scores.mean())
```

Tempo Necessário: 0 minutos and 8.36 segundos.

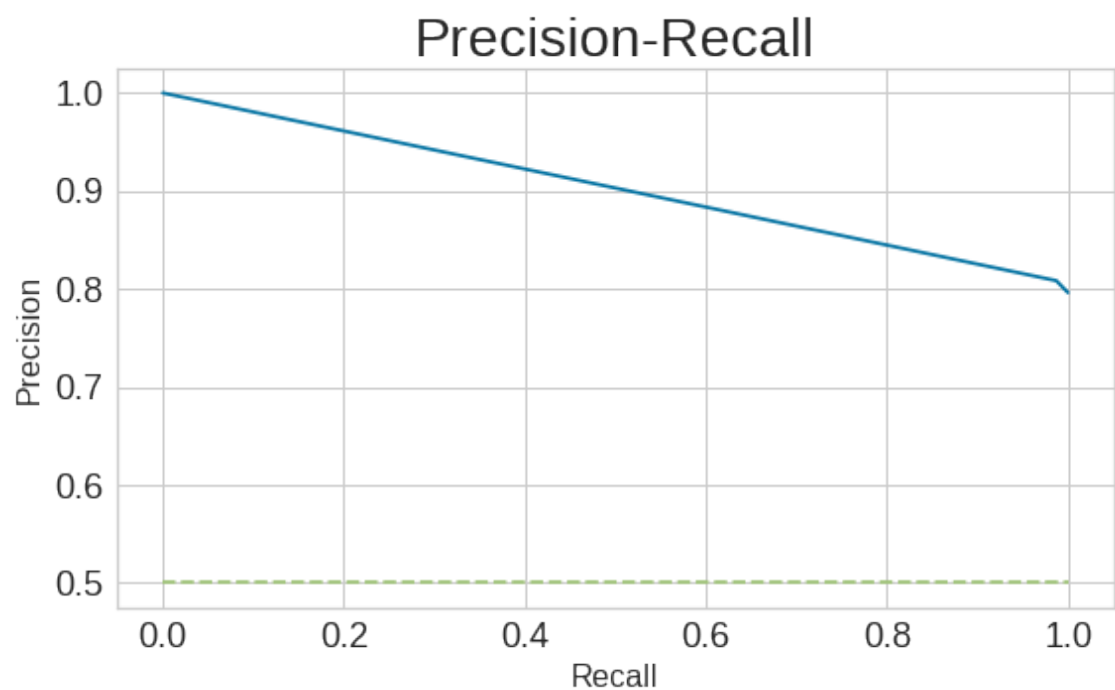
<Figure size 432x288 with 0 Axes>



<Figure size 432x288 with 0 Axes>



<Figure size 432x288 with 0 Axes>



Classification Report

	precision	recall	f1-score	support
0	0.63	0.09	0.15	13442
1	0.81	0.99	0.89	52553
accuracy			0.80	65995
macro avg	0.72	0.54	0.52	65995
weighted avg	0.77	0.80	0.74	65995

Acurácia do Modelo

Acurácia do Modelo Cross Validation

0.8002945638030585

```
[12]: plt.figure()
      plot_learning_curves(X_train, Y_train.values.ravel(), X_test, Y_test.values.
      ↪ravel(), xgb_clf, print_model=False, style='ggplot')
      plt.show()
```

